
TerraBit: Low-Bit Quantization of Geospatial Embedding Lakes

Isaac Corley
Independent Researcher
isaac.corley@proton.me

Abstract

Geospatial foundation models produce planet-scale *embedding lakes*—large, centralized stores of precomputed vectors analogous to data lakes—whose storage and retrieval costs dominate deployment budgets. Recent systems distribute global embedding fields at `int8` precision, but the quality–cost trade-off below 8 bits remains unexplored for remote sensing embeddings. We present **TerraBit**, a post-hoc quantization study spanning `float16` through 1-bit binary on 49,785,116 real 1024D embeddings (*183 GiB*), evaluated with label-free neighborhood-preservation metrics grounded in intrinsic-dimension analysis. `Int4` retains 91% kNN recall at $6.3\times$ compression; binary achieves $16.5\times$ compression with 65% recall, sufficient for coarse-stage candidate generation.

1 Introduction

Geospatial foundation models [Mendieta et al., 2023, Jakubik et al., 2024, Clay Foundation, 2025] produce high-dimensional embeddings for retrieval, change detection, deduplication, and spatial indexing, forming *embedding lakes*—large, centralized vector stores analogous to data lakes. At planet scale, storage and memory bandwidth—not model accuracy—become the deployment bottleneck. A corpus of 50M embeddings at $d=1024$ in `float32` occupies *183 GiB*; halving per-vector cost saves nearly *100 GiB* with no retraining.

Two recent systems highlight this tension. AlphaEarth releases global embedding fields cast to `int8` via post-hoc calibration [Brown et al., 2025]. TESSERA applies quantization-aware training (QAT) to a learned 128D bottleneck and distributes the result at `int8` [Feng et al., 2025]. Both treat `int8` as the practical floor. **TerraBit** asks: *how far below int8 can we push post-hoc quantization while preserving nearest-neighbor structure?*

Figure 1 previews the answer. We sweep seven bit-widths from `float16` to 1-bit binary on a real 49.8M-embedding corpus and measure neighborhood preservation with label-free metrics. The frontier is smooth: `int4` retains 91% kNN recall at $6.3\times$ compression, while binary reaches $16.5\times$ with recall sufficient for shortlist retrieval [Wang et al., 2018].

Our contributions are: (1) an intrinsic-dimension-guided evaluation framework combining ID diagnostics [Levina and Bickel, 2004, Facco et al., 2017] with label-free retrieval fidelity metrics; (2) a systematic post-hoc quantization study (`float16`, `fp8`, `int8`, `int4`, `int3`, `int2`, binary) on 49.8M real geospatial embeddings; (3) a schema-preserving parallel Parquet transformation pipeline that processes the full corpus in 45 min at 18,359 rows/s; and (4) a retrieval benchmark showing $>1,300$ QPS on GPU with recall preserved across hardware.




Figure 1: **Quality–compression Pareto frontier.** kNN recall@10 (cosine and Euclidean) vs. compression ratio on 49,785,116 geospatial embeddings. Int8 (dashed line) is the current practical baseline [Brown et al., 2025, Feng et al., 2025]. Int4 preserves 91% recall at 6.3 $\times$; binary reaches 16.5 $\times$ with useful shortlist-level recall.

2 Related Work

Geospatial foundation models. Recent work trains large vision encoders on satellite imagery and releases pretrained embeddings as data products [Mendieta et al., 2023, Jakubik et al., 2024, Clay Foundation, 2025]. AlphaEarth [Brown et al., 2025] and TESSERA [Feng et al., 2025] both distribute global int8 embedding fields, recognizing that storage cost matters as much as model quality at deployment scale.


Vector quantization and retrieval. Product quantization [Jégou et al., 2011], learned anisotropic quantization [Guo et al., 2020], and locality-sensitive hashing [Indyk and Motwani, 1998] compress embeddings for approximate nearest-neighbor search. FAISS [Johnson et al., 2021] implements these at billion scale. LLM.int8() [Dettmers et al., 2022] demonstrates that aggressive quantization preserves model quality in language; Matryoshka representations [Kusupati et al., 2022] show that embedding dimensions can be truncated with graceful degradation.

Positioning. Prior vector-compression work targets web-scale text/image embeddings or model-weight quantization. **TerraBit** focuses on the less-studied regime of post-hoc quantization of geospatial embedding lakes, where spatial metadata must be preserved alongside vectors, and label-free evaluation is the practical reality.

3 Method

Problem setup. Given a corpus $X = \{x_i\}_{i=1}^N$, $x_i \in \mathbb{R}^d$, stored in partitioned Parquet files with spatial/temporal metadata columns, we apply a quantizer $\hat{x}_i = \mathcal{Q}_m(x_i)$ under method m and evaluate each method by storage footprint, neighborhood preservation, and end-to-end throughput. Our corpus comprises 49,785,116 embeddings ($d=1024$) produced by the Clay v1.5 foundation model [Clay Foundation, 2025], a ViT trained via masked autoencoding on multi-sensor satellite imagery.

Intrinsic-dimension diagnostics. Before compression, we estimate the intrinsic dimension (ID) of the embedding manifold using MLE [Levina and Bickel, 2004], TwoNN [Facco et al., 2017], and local PCA participation ratio across preprocessing variants (raw, centered, L2-normalized). ID estimates



(a) ID estimates. All estimators converge to $\hat{d} \approx 13\text{--}17$, far below $d=1024$. (b) kNN recall (cosine) across $k \in \{10, 25, 50\}$. Recall degrades gracefully as bit-width decreases.

Figure 2: **Intrinsic dimension and neighborhood preservation.** (a) The embedding manifold has $\hat{d} \approx 13\text{--}17$, motivating aggressive quantization of redundant dimensions. (b) Neighborhood preservation declines smoothly from int8 to binary, enabling application-specific operating-point selection.

guide expectations for viable compression depth: if the data manifold has effective dimension $\hat{d} \ll d$, aggressive quantization should not collapse local neighborhood structure.

Quantization methods. We evaluate seven post-hoc methods spanning 16 bits to 1 bit: `float16` (IEEE half-precision truncation); `fp8 E4M3` (8-bit floating point with 4-bit exponent); `int $b$` for $b \in \{8, 4, 3, 2\}$ (per-dimension affine quantization with stored scale/zero-point vectors, followed by fixed-width bit packing for $b < 8$); and `binary` (sign-bit encoding, $\text{sign}(x_i) \in \{0, 1\}^d$, packed to $d/8$ bytes, supporting Hamming-distance search [Wang et al., 2018]).

Evaluation metrics. We report label-free, retrieval-oriented metrics: **kNN recall@ k** ($k \in \{10, 25, 50\}$) under cosine and Euclidean distance, measuring the fraction of true k -nearest neighbors preserved after quantization; **cosine correlation**, the Spearman rank correlation of pairwise cosine similarities (original vs. quantized); and **reconstruction cosine similarity** and MSE (when dequantization is available).

Transformation pipeline. The pipeline processes Hive-partitioned Parquet files in parallel. For each file, it reads the embedding column, applies quantizers, replaces only that column, and writes output preserving directory structure, non-embedding columns, and schema. Quantization parameters (scale, zero-point, method, bit-width) are stored in Parquet footer metadata for downstream decompression.

4 Experiments

Dataset. We evaluate on 49,785,116 embeddings ($d=1024$, `float32`) from a geospatial foundation model, stored across 1,021 Parquet files totaling *183 GiB*. Files are Hive-partitioned by level-2 geohash (511 cells), year, and month. Each record carries spatial metadata (geohash, bounding box, WKB geometry), temporal metadata, and provenance identifiers alongside the embedding vector.

Protocol. We evaluate along two axes: quality and systems efficiency. For quality, we compute label-free fidelity metrics on a 20,000-row reservoir sample (reconstruction metrics are exact over all 49,785,116 rows). For systems, we run the full-corpus transformation over all 7 methods with 8 parallel workers, reporting wall-clock time and I/O throughput.

ID estimation on a 10,000-sample subset yields MLE $\hat{d} = 17.0$, TwoNN $\hat{d} = 12.6$, and LPCA $\hat{d} = 17.0$ (Figure 2a), confirming substantial structural redundancy relative to ambient dimension $d=1024$.

Quality–compression frontier. Table 1 summarizes quality across all methods. Int8 yields near-lossless retrieval (recall@10 = 0.99) at $3.5\times$ compression, matching the precision used by AlphaEarth and TESSERA [Brown et al., 2025, Feng et al., 2025]. Int4 provides a strong balanced operating

Table 1: **Quality–compression results on 49,785,116 embeddings.** Reconstruction cosine is exact over all rows; kNN recall and cosine correlation are sampled from a 20,000-row reservoir. Best compression in **bold**.

Method	Ratio ($\times$)	Recon. Cos $\uparrow$	C@10 $\uparrow$	E@10 $\uparrow$	Cos Corr $\uparrow$
float16	1.86	1.00	1.00	1.00	1.00
fp8	3.47	1.00	0.95	0.95	1.00
int8	3.45	1.00	0.99	0.99	1.00
int4	6.29	1.00	0.91	0.91	1.00
int3	7.94	0.98	0.79	0.79	0.99
int2	11.41	0.91	0.54	0.53	0.95
binary	16.53	0.49	0.65	0.66	0.79




Figure 3: **Storage footprint per quantization method.** Original float32 corpus occupies *183 GiB*. Binary reduces this to *11 GiB* ($16.5\times$); int4 to *29 GiB* ($6.3\times$).

point: $6.3\times$ compression with 91% kNN recall. Binary achieves the highest compression ($16.5\times$) with recall@10 of 0.65—suitable as a coarse first-pass filter before re-ranking with higher-fidelity representations [Wang et al., 2018, Jégou et al., 2011].

Storage and throughput. Figure 3 shows the per-method storage footprint. Binary reduces the *183 GiB* corpus to *11 GiB*; int4 to *29 GiB*. The full pipeline (all 7 methods, 8 workers) processes the entire corpus in 45.2 min at 18,359 rows/s and 107 MiB/s output throughput.

Discussion. Three findings emerge. First, ID estimates of 13–20 (vs. ambient $d=1024$) explain why aggressive quantization does not immediately collapse neighborhood structure: most per-dimension variance is redundant for local geometry. Second, the quality decline from int8 to int2 is gradual (Figure 2b), enabling practitioners to select operating points based on their recall requirements versus storage budget. Third, binary quantization is not quality-preserving in a strict sense, but its $16.5\times$ compression and compatibility with Hamming-distance search make it attractive for first-pass candidate generation in multi-stage retrieval pipelines [Wang et al., 2018].

Retrieval benchmark. To validate that quantization-induced compression translates to practical search speedups, we benchmark brute-force kNN retrieval on a 1M-vector subset (1,000 queries, $k=50$) using FAISS [Johnson et al., 2021] on CPU and PyTorch cdist on a single NVIDIA RTX 3090 GPU. Each quantized corpus is dequantized to float32 before indexing, isolating the effect of quantization noise from format-specific index optimizations; binary embeddings use Hamming distance directly via packed uint8 search.

Table 2 and Figure 4 summarize throughput and recall. GPU search achieves 1,322–1,362 QPS for dequantized methods—a $22\times$ speedup over FAISS CPU (59–60 QPS)—with identical recall, confirming that quantization quality is hardware-independent. Binary Hamming search runs at 801 QPS on CPU and 164 QPS on GPU; the lower GPU throughput reflects the overhead of bit-

Table 2: **Brute-force search benchmark** on 1M embeddings (1K queries, $k=10$). QPS = queries per second. GPU is a single RTX 3090. Recall is identical across backends since quantization quality is hardware-independent. **Bold**: best throughput per row.

Method	CPU QPS	GPU QPS	Speedup	Recall@10
float16	59	1322	22.4×	1.00
fp8	59	1361	23.1×	0.95
int8	60	1349	22.5×	0.48
int4	60	1337	22.3×	0.46
int3	59	1329	22.5×	0.43
int2	59	1346	22.8×	0.37
binary	801	164	0.2×	1.00*

*Recall measured against binary GT (Hamming distance).




Figure 4: **Search throughput vs. recall** on 1M embeddings. GPU brute-force delivers >1,300 QPS for all dequantized methods. Binary Hamming search is fastest on CPU (801 QPS) due to native bitwise ops. Recall (purple diamonds, right axis) tracks quantization fidelity independently of hardware.

unpacking binary vectors into float tensors for `cdist`, a limitation of the brute-force approach that purpose-built Hamming kernels would eliminate.

Limitations. We use label-free neighborhood-preservation proxies rather than task-specific metrics (e.g., land-cover classification accuracy). kNN recall and cosine correlation are computed on a reservoir sample, not the full corpus. Our retrieval benchmark uses brute-force search; approximate nearest-neighbor indices (IVF, HNSW) may alter the quality–throughput frontier at larger scale [Johnson et al., 2021]. Future work should add labeled downstream evaluation and investigate ANN index interactions with low-bit quantization.

5 Conclusion

TerraBit provides a systematic compression study for geospatial embedding lakes. On 49.8M real embeddings, low-bit quantization yields a smooth quality–cost frontier: `int4/int8` are strong balanced choices; `binary` is viable for coarse retrieval stages. GPU-accelerated brute-force search over dequantized embeddings exceeds 1,300 QPS on a single RTX 3090, demonstrating that quantized embedding lakes are practical for real-time retrieval at scale.

References

- Christopher F. Brown, Michal R. Kazmierski, Valerie J. Pasquarella, William J. Rucklidge, Masha Samsikova, Chenhui Zhang, Evan Shelhamer, Estefania Lahera, Olivia Wiles, Simon Ilyushchenko, Noel Gorelick, Lihui Lydia Zhang, Sophia Alj, Emily Schechter, Sean Askay, Oliver Guinan, Rebecca Moore, Alexis Boukouvalas, and Pushmeet Kohli. AlphaEarth Foundations: An embedding field model for accurate and efficient global mapping from sparse label data. *arXiv preprint arXiv:2507.22291*, 2025.
- Clay Foundation. Clay foundation model v1.5. <https://clay-foundation.github.io/model/>, 2025. Apache-2.0 license.
- Tim Dettmers, Mike Lewis, Younes Belkada, and Luke Zettlemoyer. LLM.int8(): 8-bit matrix multiplication for transformers at scale. In *Advances in Neural Information Processing Systems*, 2022.
- Elena Facco, Maria d’Errico, Alex Rodriguez, and Alessandro Laio. Estimating the intrinsic dimension of datasets by a minimal neighborhood information. *Scientific Reports*, 7(1):12140, 2017.
- Zhengpeng Feng, Clement Atzberger, Sadiq Jaffer, Jovana Knezevic, Silja Sormunen, Robin Young, Madeline C. Lisaius, Markus Immitzer, Toby Jackson, James Ball, David A. Coomes, Anil Madhavapeddy, Andrew Blake, and Srinivasan Keshav. TESSERA: Temporal embeddings of surface spectra for earth representation and analysis. *arXiv preprint arXiv:2506.20380*, 2025.
- Ruiqi Guo, Philip Sun, Erik Lindgren, Quan Geng, David Simcha, Felix Chern, and Sanjiv Kumar. Accelerating large-scale inference with anisotropic vector quantization. In *International Conference on Machine Learning*, pages 3887–3896, 2020.
- Piotr Indyk and Rajeev Motwani. Approximate nearest neighbors: Towards removing the curse of dimensionality. In *Proceedings of the Thirtieth Annual ACM Symposium on Theory of Computing*, pages 604–613, 1998.
- Johannes Jakubik, Sujit Roy, C. E. Phillips, Paolo Fraccaro, Denys Godwin, Birgit Zadrozny, Daniela Szwarcman, Carlos Gomes, Gabby Nyirjesy, Blair Edwards, Manil Maskey, et al. Foundation models for generalist geospatial artificial intelligence. *arXiv preprint arXiv:2310.18660*, 2024.
- Hervé Jégou, Matthijs Douze, and Cordelia Schmid. Product quantization for nearest neighbor search. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 33(1):117–128, 2011.
- Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search with GPUs. *IEEE Transactions on Big Data*, 7(3):535–547, 2021.
- Aditya Kusupati, Gantavya Bhatt, Aniket Rege, Matthew Wallingford, Aditya Sber, Prateek Jain, Sham Kakade, and Ali Farhadi. Matryoshka representation learning. In *Advances in Neural Information Processing Systems*, 2022.
- Elizaveta Levina and Peter J. Bickel. Maximum likelihood estimation of intrinsic dimension. *Advances in Neural Information Processing Systems*, 17, 2004.
- Matias Mendieta, Boran Han, Xingjian Shi, Yi Zhu, and Chen Chen. Towards geospatial foundation models via continual pretraining. *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 16806–16816, 2023.
- Jingdong Wang, Ting Zhang, Jingkuan Song, Nicu Sebe, and Heng Tao Shen. A survey on learning to hash. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 40(4):769–790, 2018.